

A Technical Comparison of Multi-agent Systems Tools

Hamid Nemati

Supervised by: Dr. Dokyun Lee

Boston University, CDS, BITLAB

September, 2024

Abstract

Large Language Models (LLMs) have limitations, but when combined in multi-agent systems, they can address complex problems more effectively than a single LLM. This report examines Multi-Agent Systems (MAS), detailing their foundational principles, potential applications, and the various frameworks developed to support them. Among these frameworks, AutoGen, Taskweaver, and LangGraph stand out as the most stable and promising for advancing MAS capabilities, offering robust solutions for complex, multi-agent interactions.

1. Introduction

Multi-agent Large Language Models (LLMs) represent an innovative approach in artificial intelligence, where multiple specialized agents collaborate to solve complex tasks. Unlike traditional single-agent models, these systems leverage the unique strengths of each agent, allowing them to outperform in scenarios that require diverse skills and viewpoints. The core of this teamwork is the ability of agents to work either independently or in tandem, depending on the task at hand, with human oversight to ensure accuracy and efficiency.

1.1. What are Multi-Agent LLMs?

Multi-agent LLMs consist of language models that are teamed up, each taking on a specific role that aligns with its strengths. These systems excel in areas where single-agent models may struggle, particularly in tasks requiring multifaceted approaches or real-world applications. The collaboration among agents is what sets multi-agent LLMs apart, allowing them to effectively pull together the strengths of different specialized agents to tackle complex problems.

These agents operate either autonomously or in collaboration, depending on the task requirements. While they largely function independently, human supervision is essential for overseeing their decisions and validating their outputs. Equipped with powerful language models, agents utilize various tools to perform tasks such as web searches or document processing.

1.2. How Multi-Agent LLMs Work

A typical multi-agent LLM system workflow begins with a user providing a high-level task or query. The system then decomposes this task into smaller, manageable subtasks, assigning them to the appropriate specialized agents based on their capabilities. Each agent utilizes its LLM to reason, plan, and execute the task using its available tools and memory. Inter-agent communication is key, as agents share

information as needed to complete interdependent subtasks, culminating in a final output assembled from the collective efforts of all agents.

1.3. Why Multi-Agent LLMs Are Superior

Multi-agent LLMs offer several advantages over single-agent systems:

1. Accuracy and Hallucination Reduction: Single-agent LLMs are prone to hallucinations—generating plausible but incorrect information, which is a critical issue in fields like medicine or law. Multi-agent systems mitigate this problem by enabling agents to cross-check each other’s work, significantly reducing errors and enhancing reliability. Research indicates that employing multiple agents can improve accuracy and make these systems particularly valuable in critical settings.

2. Extended Context Handling: Single-agent LLMs are limited by their context windows, which restrict the amount of text they can process at one time. Multi-agent systems overcome this limitation by distributing the workload among several agents, each focusing on different segments of text. This collaboration allows for a more comprehensive and continuous understanding of extended documents or long-term conversations.

3. Efficiency and Multitasking: Single-agent LLMs process tasks sequentially, leading to delays in scenarios requiring quick responses to multiple queries. In contrast, multi-agent systems enhance efficiency through parallel processing, where multiple agents handle different tasks simultaneously. This approach reduces response times and boosts productivity, making it ideal for business environments where promptness is crucial.

4. Collaborative Capabilities: Multi-agent systems excel in tasks requiring teamwork, bringing together the diverse skills and expertise of various agents. This collaboration is particularly beneficial in areas like scientific research or strategic planning, where combining different perspectives leads to better outcomes.

While single-agent systems are effective for cognitive tasks that can be handled independently, multi-agent systems are better suited for complex, dynamic tasks. Each

agent in a multi-agent system brings unique problem-solving methods and collaborates with others to achieve shared goals.

1.4. Building Blocks of Multi-Agent Systems

1. Large Language Models (LLMs): LLMs serve as the intellectual core of modern MAS, providing the reasoning and decision-making capabilities needed for sophisticated tasks.

2. Agents: Agents are autonomous entities designed to perform specific tasks, make decisions, and work collaboratively towards common objectives. Using LLMs as their reasoning mechanisms, these agents can adapt to changing conditions and interact dynamically within the system. Their autonomy allows them to efficiently handle a variety of tasks and challenges.

3. Environment: The environment refers to the external context or space in which agents operate. It includes all the external factors and resources that agents interact with, such as data sources, users, and other agents. The environment provides the stimuli that agents respond to and adapt within.

4. Tools and Resources: Tools are specialized functions or skills that agents use to accomplish their tasks, ranging from simple actions like data retrieval to complex operations such as detailed analyses.

5. Processes or Flows: Processes define how tasks are organized and executed within a MAS. They ensure that tasks are distributed efficiently and aligned with overall objectives, whether through inter-agent collaboration or intra-agent interaction with tools and outputs.

6. Communication Protocols: Communication is essential for collaboration in MAS. Protocols define the rules and languages through which agents exchange information, negotiate, and coordinate their actions. Effective communication protocols are necessary for ensuring that agents can work together seamlessly, especially in complex, dynamic environments.

7. Coordination Mechanisms: These mechanisms are essential for managing how agents work together towards common goals. Coordination can involve task alloca-

tion, conflict resolution, and synchronization of actions, ensuring that the overall system operates efficiently and effectively.

8. Learning and Adaptation: MAS often includes mechanisms for learning from experience and adapting to changes in the environment or objectives. This could involve reinforcement learning, supervised learning, or unsupervised methods that allow agents to improve their performance over time.

1.5. How Multi-Agent Systems Can Help

1. Enhanced Efficiency: MAS automate complex workflows, reducing human error and increasing productivity in business environments.

2. Data-Driven Decisions: MAS provide real-time data analysis and insights, enabling quick decision-making and responsiveness to changing conditions.

3. Personalized Customer Service: Intelligent agents within MAS offer 24/7 customer support, handling inquiries, resolving issues, and guiding customers through various processes.

4. Resilient Supply Chains: MAS enhance supply chain resilience by enabling quick adaptation to disruptions, monitoring activities, detecting anomalies, and implementing contingency plans.

5. Collaborative Innovation: MAS accelerate collaborative innovation by bringing together diverse expertise and perspectives, coordinating efforts, and contributing to problem-solving in a cohesive manner.

1.6. Key Characteristics of Agents

Agents in MAS exhibit several key characteristics:

- **Autonomy:** Operates independently, making decisions.
- **Goal-directed Behavior:** Focuses on achieving goals through perception, reasoning, and action.
- **Social Ability:** Communicates and interacts with others.

- **Reactivity:** Responds to changes in the environment.
- **Proactivity:** Takes initiative, exhibiting goal-directed behavior.
- **Adaptability:** Learns from experience to modify behavior.
- **Perception:** Gathers information from the environment.
- **Action:** Uses tools to perform actions that affect the environment.
- **Learning:** Improves performance over time through experience.
- **Reasoning:** Processes information and makes decisions based on logic.
- **Planning:** Breaks down complex tasks into smaller ones and outlines a plan for execution.
- **Communication:** Exchanges information effectively with the environment and other agents.

1.7. Multi-Agent Architectures

MAS architectures can be categorized as:

- **Sequential:** Agents perform tasks in a predetermined order, where each step depends on the completion of the previous one.
- **Hierarchical:** Agents follow a structured hierarchy, with higher-level agents overseeing and coordinating the activities of lower-level agents.
- **Asynchronous:** Agents operate independently, handling tasks as they arise without adhering to a fixed sequence, allowing for flexibility and parallel processing.

There are three types of **agent architectures**:

- **Reactive Agents:** Operate on a stimulus-response basis without maintaining an internal state. Suitable for dynamic environments but limited in handling complex tasks.
- **Deliberative Agents:** Use an internal model of the environment to plan actions, capable of reasoning and long-term planning, but require higher computational resources.
- **Hybrid Agents:** Combine elements of both reactive and deliberative architectures, balancing responsiveness and strategic capability.

1.8. Agent Communication

Effective communication among agents is crucial in MAS:

- **Coordination:** Ensures agents' actions are aligned toward the system's objectives.
- **Cooperation:** Enables agents to share resources and collaborate on tasks.
- **Negotiation:** Helps resolve conflicts in competitive or resource-constrained environments.
- **Information Sharing:** Enhances decision-making by sharing information about the environment and agents' internal states.

Agents communicate through various methods:

- **Message Passing:** Agents exchange messages using a predefined protocol, which can be synchronous or asynchronous.
- **Broadcasting:** Useful for quickly disseminating information, though it can lead to network congestion.
- **Direct Communication:** Involves direct exchanges between specific agents, reducing communication overhead and increasing efficiency.

- **Middleware:** Middleware solutions, such as agent platforms and frameworks like CrewAI, LangGraph, and AutoGen, provide essential services that facilitate communication and coordination among agents. (In Multi-Agent Systems (MAS), middleware refers to the software infrastructure that supports seamless agent operation across various systems. It typically includes: **Agent Platforms**, which offer a runtime environment for agents and provide essential services such as messaging, directory management, and lifecycle management to ensure efficient operation within the system; **Agent Communication Languages (ACL)**, such as FIPA ACL, which enable agents to exchange messages in a structured and coherent manner, promoting compatibility across platforms; and **Interoperability Protocols**, which allow agents with different architectures or running on different platforms to communicate and collaborate effectively, enhancing the flexibility and scalability of MAS.)

1.9. Challenges in MAS Implementation

Despite the advantages of MAS, challenges exist in their implementation:

- **Scalability:** Managing a large number of agents, each with complex behaviors, can lead to increased computational and communication overhead.
- **Security:** Ensuring that agents behave as expected and do not become compromised is critical, particularly in open systems where malicious agents might be introduced.
- **Coordination and Cooperation:** As the number of agents increases, maintaining effective coordination and cooperation becomes more difficult, potentially leading to inefficiencies or conflicts.
- **Real-Time Decision Making:** In dynamic environments, agents must make decisions quickly, which can be challenging when balancing the need for fast responses with the need for accurate and well-reasoned decisions.

2. Detailed Comparison

When new frameworks emerge, they first need to prove their capabilities. However, as they evolve, new concepts often force these frameworks to grow more complex, rather than simplifying. This can make things more cumbersome over time. Take Langchain, for instance—it was promising at first, but later required the addition of LangGraph due to its inability to handle Directed Acyclic Graph (DAG) structures effectively. Even with LangGraph, it struggles with highly interconnected tasks that demand extensive agent collaboration beyond basic graph structures.

Then came AutoGen and CrewAI, aimed at simplifying the workflow, yet controlling the execution plan remains challenging—even with the use of contexts. AutoGen allows developers to review the “inner monologue” of agent workflows, offering some transparency, but CrewAI functions more like a black box, giving us limited control over the process. Microsoft, dissatisfied with AutoGen’s limitations, introduced TaskWeaver. As a result, a growing number of these frameworks flood the market, with new ones appearing almost daily.

In this section, we will discuss the most significant frameworks in this evolving landscape. Let’s first categorize them into three groups and then we talk about each group in details: (There isn’t a formalized categorization of Multi-Agent Systems (MAS) tools into groups across academic literature and this categorization is made by the author of this report.)

Static Path MAS: This group could include tools like LangChain, LangGraph, and Langroid. These frameworks follow predefined workflows or chains where the user defines the entire sequence of interactions, with limited decision-making flexibility. They excel in situations where the path of execution is clear and relatively deterministic. Each step in the chain is known ahead of time, and the agents follow the logic provided by the developer.

Dynamic Path MAS: Tools like AutoGen and CrewAI fall into this category. Here, the path of agent interactions is more flexible and determined dynamically, often by a “planner” agent. These frameworks allow for on-the-fly decision-making where agents assess tasks in real time, adjusting their behavior based on evolving contexts

or inputs. The execution path is not rigidly predefined, offering more adaptability to unpredictable environments.

Domain-Specific MAS: This category would encompass tools like OpenDevin and MetaGPT, which are designed for specialized domains (such as software engineering). These systems come with predefined agent roles tailored to particular tasks. The focus is on solving specific problems, like automating the software development process, where the roles and interactions between agents are already well-structured for that domain.

2.1. Static Path MAS

2.1.1. LangChain

LangChain is a framework designed for building applications powered by Large Language Models (LLMs), with a focus on simplifying the creation, customization, and deployment of these systems. While it is not explicitly a multi-agent system framework, LangChain enables the orchestration of multiple agents working together on complex tasks, allowing for task-solving workflows.

LangChain is well-known for its ReAct framework, which structures task-solving through few-shot steps involving reasoning, actions, and feedback from the environment. However, several limitations of LangChain have been noted:

- **Complexity:** The framework can be overly complex, especially when modifying simple prompts. Such changes often require navigating deep into the code, making it cumbersome for developers.
- **Sparse Documentation:** While LangChain’s documentation is sufficient for basic use cases, it lacks depth for advanced tasks like custom agent development. Max Woolf highlighted this in his blog, remarking on the challenges of understanding LangChain’s structure.
- **Inconsistent Output:** Achieving consistent output in formats like JSON is challenging, particularly when using models beyond OpenAI’s. LangChain doesn’t fully address these challenges, leading to inconsistencies.

- **Difficult Orchestration:** Implementing complex orchestration algorithms requires a deep dive into the source code, making it less user-friendly for users wanting more sophisticated multi-agent coordination.

Despite these shortcomings, LangChain has contributed significantly to making LLMs accessible to a wider audience, particularly through its popularization of Retrieval-Augmented Generation (RAG). Moreover, it pairs well with Llama-Index, a library that simplifies the loading and management of data for LLM applications. However, the discussion of Llama-Index will be reserved for a future article.

2.1.2. LangGraph

LangGraph is a specialized framework designed for building advanced multi-agent systems with cyclical, stateful workflows. It allows for more intricate agent orchestration compared to LangChain, offering developers fine-grained control over how agents interact, communicate, and perform iterative tasks.

The framework is particularly useful for scenarios requiring cyclical computation and agent collaboration, enabling agents to engage in dynamic decision-making and multi-step reasoning. This is achieved through a graph-based structure that allows developers to define and execute sophisticated workflows.

Key features of LangGraph include:

- **Stateful Multi-Agent Applications:** LangGraph supports applications where multiple agents maintain their state over long interactions, making it ideal for systems requiring continuous collaboration.
- **Cyclical Computation:** LangGraph allows for cyclical processes, such as simulations and iterative reasoning, where agents repeatedly interact to refine their outputs.
- **Fine-Grained Control:** The framework provides developers with more control over agent behavior and interactions compared to LangChain, making it suitable for complex, custom-built systems.

While LangChain simplifies multi-agent orchestration for beginners and intermediate developers, LangGraph caters to those needing more detailed control and flexibility. LangGraph is particularly advantageous in applications that require robust, cyclical workflows where agents must repeatedly collaborate and review each other's work. While AutoGen excels in conversational engagement, CrewAI provides structured flexibility, LangGraph allows for detailed control over agent interactions through graphical representation.

2.1.3. Langroid

Langroid is a code-first framework for managing complex multi-agent workflows, with a focus on flexibility and direct control over agent interactions. It is designed to enable more advanced multi-agent coordination by allowing developers to define modular and graph-like workflows that agents follow, offering a higher degree of customization than LangChain or LangGraph.

Langroid builds upon the limitations of existing frameworks by offering better support for highly interconnected tasks. Its key strength lies in enabling more sophisticated agent collaboration, particularly in scenarios where agents must maintain memory and context across interactions.

Key characteristics of Langroid include:

- **Dynamic Multi-Agent Workflows:** Langroid enables flexible, modular workflows where agents collaborate dynamically, adjusting their actions based on the evolving context of the task.
- **Task Delegation and Memory:** Langroid supports agents that can delegate tasks to other agents while maintaining context, making it suitable for complex orchestration scenarios.
- **Direct Control:** Unlike LangChain, which abstracts much of the orchestration process, Langroid provides a more hands-on approach, giving developers greater control over agent behavior and task flow.

Langroid distinguishes itself from LangChain and LangGraph by providing a more flexible framework for managing highly interconnected, dynamic workflows. While LangChain is simpler to use for single-agent or NLP-centric tasks, and LangGraph is more structured for cyclical workflows, Langroid was created to fill the need for advanced multi-agent coordination in larger, more complex systems.

2.2. Dynamic Path MAS

The balance between exploration and exploitation is central to decision-making in multi-agent systems (MAS), particularly in contexts where agents must operate autonomously, adapt to dynamic environments, and achieve optimal collective outcomes. In this subsection we review these kind of MAS.

2.2.1. AutoGen

In September 2023, Microsoft released AutoGen – a flexible and open-source Python-based framework for defining, configuring, and composing AI agents to drive multi-agent applications with minimal coding. Its low-code interface, AutoGen Studio, provides a library of pre-defined agents that can be composed into teams (workflows). These agents are highly customizable, supporting integration with LLMs, tools, and human feedback, making them suitable for a wide range of applications like personal finance management and learning assistance.

AutoGen Studio builds on AutoGen, providing a user-friendly interface for creating and customizing agents with minimal coding. It allows developers to quickly build, test, deploy, and share agents and workflows within the community. AutoGen Studio features a “Build” section where users can select from a library of pre-defined agents and compose them into teams (workflows). Developers can test these workflows on various tasks, review resulting artifacts (such as images, code, and documents), and examine the “inner monologue” of agent workflows to understand their task management. Additionally, AutoGen Studio offers profiling information, including costs associated with the run (such as the number of turns and tokens) and agent actions

(such as tool calls and code execution outcomes). **Unlike CrewAI, AutoGen Studio provides more transparency and control over the process.**

While AutoGen is not as popular as LangChain, it is a newer and promising contender. My experience with AutoGen, both professionally and personally, has been positive. It excels in applications requiring dynamic agent cooperation but may be less suited for highly controlled environments due to its inherent flexibility. Currently, AutoGen is a top choice for LLM orchestration and tool-augmented applications.

Pros

- **Well-established:** AutoGen benefits from an active community, providing valuable support and collaboration opportunities for developers.
- **Customizable Agents:** The framework offers highly customizable agents that can integrate with LLMs, tools, and human feedback, enabling flexible task execution.
- **Easy Configuration:** Agent configuration is straightforward and user-friendly.
- **Versatile Model Support:** AutoGen supports various models including OpenAI, OpenRouter, Azure OpenAI, and even custom models hosted on Azure.
- **Simple System Prompts:** Interacting with system prompts is effortless.
- **Rich Examples:** Despite somewhat lacking documentation, the extensive examples available in the GitHub repository make AutoGen a viable choice for newcomers.
- **Simplified Complex Orchestration:** The GroupChat class and data flow facilitate the creation of complex, custom orchestration systems.
- **Tool Augmentation:** Integrating additional tools is straightforward.
- **Modularity:** AutoGen's modular nature allows it to work seamlessly with various libraries, such as Llama-Index for tool augmentation.

- **Production-Ready:** Compared to alternatives like LangChain, AutoGen offers faster performance and more effective chain-of-thought processing.
- **Open-Source:** As an open-source framework with a growing community and dedicated contributors, AutoGen continues to evolve and improve.

Cons

- **Steep Learning Curve:** The framework’s complexity may present a challenge for new users, requiring a significant learning curve to master its features.
- **Less Structured:** Compared to other frameworks, AutoGen may appear less structured, potentially affecting ease of implementation.

AutoGen FAQ

- Can Autogen agents communicate with external systems or APIs? Yes.
- Is Autogen limited to Python code generation and execution? No.
- Can I integrate Autogen into my existing workflow or project? Yes.
- Can Autogen agents work offline or in a restricted network environment? Not really.
- Can Autogen agents be trained or customized for specific domains or industries? Yes.
- What are the potential limitations or challenges of using Autogen? training data
- availability, model customization, and computational resources.

2.2.2. TaskWeaver

Microsoft built TaskWeaver to address limitations in AutoGen and meet more specific demands in multi-agent systems and workflow automation. TaskWeaver was introduced to enhance the orchestration of multi-agent tasks by providing more advanced handling of intricate, interconnected processes and ensuring more efficient delegation and coordination among agents. It aims to offer greater control over execution plans, which was a challenge in AutoGen. **Currently, TaskWeaver is still in active development, with changes being made almost every month.**

Microsoft refers to TaskWeaver as a "code-first" agent framework. This means that TaskWeaver is designed for developers who prioritize direct coding and programmatic control when building and orchestrating agents. In a "code-first" approach, the system is primarily interacted with through writing code, rather than using a visual interface or no-code/low-code solutions. Code-first means:

- **Developer-Centric Control:** Developers use programming languages (likely Python or other supported languages) to define how agents are created, configured, and interact with each other. This contrasts with a "no-code" or "low-code" environment, where users rely on pre-built templates, drag-and-drop interfaces, or visual tools to build workflows.
- **Fine-Grained Customization:** A code-first approach allows developers more flexibility and granularity in customizing agent behaviors, decision-making processes, and workflow orchestration. It gives them control over logic, data flows, and the integration of external tools or systems.
- **Scalability and Precision:** For complex, large-scale projects that require tight control over agent interactions and task execution (like AI research, enterprise-level applications, or specialized domains), a code-first framework like TaskWeaver can handle more sophisticated use cases. It ensures that developers can optimize performance, debug issues at a low level, and tailor the agent system precisely to their needs.

Among its most important characteristics of TaskWeaver are:

- **Modular Design:** TaskWeaver decomposes tasks into smaller, manageable subtasks that can be executed independently by different LLMs or reasoning components. This modularity simplifies handling complex tasks and allows subtask reuse across different workflows.
- **Declarative Specification:** Tasks are defined in a declarative language that focuses on the desired outcome rather than the specific method to achieve it. This enables TaskWeaver to generate execution plans dynamically based on available resources and current conditions.
- **Contextual Reasoning:** TaskWeaver incorporates contextual information from prior subtasks, external APIs, and real-world data, allowing it to make tailored decisions based on the specific circumstances of each task.
- **Fault Tolerance:** TaskWeaver is built to handle errors and failures during execution. It can retry failed subtasks, revert to previous states, or explore alternative methods to complete tasks successfully.

Compared to other frameworks like AutoGen Studio or CrewAI, which provide user-friendly or low-code interfaces, TaskWeaver is tailored to developers who prefer greater control over agent interactions and task execution.

AutoGen is a flexible, open-source framework designed for configuring and composing AI agents with minimal coding. It is well-suited for a wide range of applications but struggles with more complex workflows and seamless agent interaction. AutoGen uses a **template-based** approach for code generation, allowing developers to customize predefined templates for different use cases.

In contrast, **TaskWeaver** is a code-first framework designed to offer precise control and advanced orchestration of multi-agent tasks. Its strengths lie in its modular design, declarative task specifications, and contextual reasoning, making it ideal for large-scale and complex projects. TaskWeaver builds LLM-powered autonomous agents by converting user requests into executable code, treating user-defined plugins as callable functions. It supports rich data structures, flexible plugin usage, and

dynamic plugin selection, leveraging LLMs for sophisticated logic. TaskWeaver also integrates domain-specific knowledge and ensures secure code execution.

2.2.3. Crew AI

CrewAI is an open-source multi-agent orchestration framework created by João Moura, a Research Associate at The University of Edinburgh. This Python-based framework leverages artificial intelligence (AI) collaboration by orchestrating autonomous AI agents that work together as a cohesive unit, or "crew," to complete tasks. The goal of CrewAI is to automate multi-agent workflows effectively.

CrewAI encourages users to think of agents as members of a team, where each agent is assigned a specific role, such as 'Data Scientist,' 'Researcher,' or 'Product Manager.' These agents collaborate autonomously to perform tasks and make decisions, communicating with each other to ensure workflow efficiency. A key aspect of CrewAI is its orchestration of processes, which define how agents work together and what tasks they will perform. These processes ensure task distribution, execution, and alignment with a predefined strategy to achieve a goal.

CrewAI supports two main process implementations:

- **Sequential Process:** Similar to a dynamic team workflow, tasks are executed in a predefined order, with the output of one task feeding into the next.
- **Hierarchical Process:** Emulates a corporate hierarchy where a manager agent, created by leveraging a manager language model, oversees task execution, allocates tasks, and reviews outputs.
- **Consensual Process (Planned):** A future process designed to allow agents to collaborate democratically on task execution, introducing a more collaborative approach to task management.

One of CrewAI's core features is the concept of tools, which are skills or functions that agents can utilize to perform actions. Tools come from both the CrewAI Toolkit and LangChain Tools, enabling everything from basic searches to complex interactions. Key characteristics of CrewAI tools include:

- **Utility:** Tools are designed for tasks like web searching, data analysis, content generation, and collaboration.
- **Integration:** Tools are seamlessly integrated into agents' workflows, enhancing their capabilities.
- **Customizability:** CrewAI allows users to develop custom tools or use existing ones, adapting to specific agent needs.
- **Error Handling:** The framework incorporates robust mechanisms to handle errors, ensuring smooth operation.
- **Caching Mechanism:** CrewAI uses intelligent caching to optimize performance by avoiding redundant operations.

CrewAI offers a **more flexible**, role-based approach compared to other frameworks like AutoGen and ChatDev. Agents in CrewAI autonomously delegate tasks, consult one another, and collaborate without constant external intervention, making it ideal for dynamic, adaptable environments.

CrewAI vs. AutoGen: AutoGen is Microsoft's open-source agentic framework using natural language processing (NLP) for conversational AI agents. Both CrewAI and AutoGen support customizable AI agents, but CrewAI offers simpler orchestration of agent interactions with customizable attributes controlling workflows. In contrast, AutoGen requires more programming setup but has built-in functionality to quickly execute LLM-generated code, something CrewAI currently lacks.

CrewAI vs. ChatDev: ChatDev is another open-source platform that supports multi-agent collaboration, including some features from CrewAI. However, ChatDev's process structure is more rigid, limiting customization and scalability. CrewAI is better suited for production environments requiring integration with third-party applications and adaptable workflows. ChatDev, on the other hand, extends itself as a browser extension, linking conversations across agents in a web browser.

2.2.4. Agency-Swarm

Agency Swarm is a comprehensive Multi-Agent System (MAS) framework designed by Arsenii Shatokhin (aka VRSEN) to facilitate the automation of AI agencies. The framework allows developers to create and oversee distributed multi-agent systems (Agencies), each composed of agents with distinct roles and capabilities. Agency Swarm emphasizes scalability, flexibility, and ease of customization, making it suitable for large-scale projects and specialized use cases.

Agency Swarm is built to handle distributed computing environments across varying scales, supporting efficient communication between agents to foster collaboration and coordination. It includes environment simulation tools, which allow developers to test and refine agent behaviors by simulating complex scenarios. The framework also offers visualization tools to monitor and understand agent interactions and system dynamics, making it easier to debug and optimize multi-agent systems.

Some of the standout features of Agency Swarm include:

- **Customizable Agent Roles:** Developers can define agent roles such as CEO, virtual assistant, or developer and customize their functionalities using the Assistants API.
- **Full Control Over Prompts:** Unlike some frameworks, Agency Swarm does not impose predefined prompts, allowing users to fully customize agent behavior.
- **Tool Creation:** Tools in Agency Swarm are created via the Instructor interface, which includes automatic type validation to prevent errors.
- **Efficient Communication:** Agents communicate using a custom "SendMessage" tool, determining communication partners based on their own descriptions.
- **State Management:** The framework manages the state of agents efficiently using a 'settings.json' file, ensuring persistent assistant states in OpenAI environments.

- **Deployability:** Agency Swarm is designed to be production-ready, making it a reliable choice for large-scale implementations.

Below we have pros and cons of Agency-Swarm: **Strengths:**

- **Scalability:** It is built to handle distributed environments, making it suitable for large-scale applications.
- **Flexibility:** Offers extensive customization, allowing developers to tailor the framework to their specific needs.
- **Simulation Support:** Provides tools to simulate environments for effective MAS development and testing.
- **Visualization:** Includes visualization tools for better understanding and debugging of MAS behaviors.

Limitations:

- **Documentation:** The documentation is reportedly lacking or incomplete, which could be challenging for new users.
- **Learning Curve:** As with many MAS frameworks, there is a steep learning curve, particularly for those new to multi-agent systems.

Unlike other frameworks, Agency Swarm does not write prompts for you but prevents hallucinations with automatic type checking and error correction with instructor which allows you to easily define communication flows. Below we have the comparison of Agency-Swarm with its rivals. **Agency Swarm vs. AutoGen:** AutoGen, Microsoft’s agentic framework, determines the next speaker using an additional model call, which mimics role-playing between agents. This can be inefficient and less customizable, as it limits the control over which agents communicate with each other. Although AutoGen recently introduced hardcoded conditions to control the next speaker, this undermines adaptability, a core feature of agentic systems. In

contrast, Agency Swarm handles communication via the "SendMessage" tool, allowing agents to determine who to communicate with based on their descriptions, while developers set the communication boundaries within the agency chart.

Agency Swarm vs. CrewAI: CrewAI introduces a "process" mechanism to control agent communication, but its reliance on LangChain—a framework developed before function-calling models were available—means it lacks type checking and error correction. This introduces the risk of system failure due to hallucinations in agent actions. CrewAI's primary advantage lies in its compatibility with open-source models, whereas Agency Swarm offers automatic type checking, error correction via the Instructor interface, and more refined communication flows between agents, making it a more reliable choice for production environments.

2.2.5. MetaGPT

MetaGPT is a multi-agent framework that in itself acts as a virtual software company. It takes a one-line requirement and outputs user stories, competitive analysis, requirements, data structures, APIs, and documents. MetaGPT includes product managers, architects, project managers, and engineers, following carefully crafted Standard Operating Procedures (SOPs). MetaGPT takes a one line requirement as input and outputs user stories / competitive analysis / requirements / data structures / APIs / documents, etc. Internally, MetaGPT includes product managers / architects / project managers / engineers. It provides the entire process of a software company along with carefully orchestrated SOPs.

Some of the standout features of MetaGPT include:

- **Collaborative Agent Roles:** MetaGPT allows the creation of agents with distinct professional roles (e.g., manager, developer, etc.) that mimic real-world job functions and responsibilities.
- **Dynamic Task Delegation:** Agents dynamically assign tasks based on their roles and responsibilities, ensuring an adaptive workflow similar to human teams.

- **Decentralized Coordination:** Agents coordinate autonomously without centralized control, making decisions based on task requirements and agent capabilities.
- **Adaptability:** MetaGPT is highly adaptable to varying workflows, making it suitable for a wide range of applications, from software development to project management.
- **Role-Playing Simulations:** By simulating team dynamics, MetaGPT allows developers to test and iterate on complex, collaborative workflows in a controlled environment.

Strengths:

- **Flexibility:** The framework is designed to simulate a wide range of real-world roles, making it adaptable to many industries.
- **Scalability:** MetaGPT can handle complex, large-scale projects by allowing agents to efficiently coordinate tasks.
- **Decentralized Coordination:** Agents can autonomously make decisions without requiring centralized control, mirroring real-world team dynamics.
- **Team Simulation:** MetaGPT's role-based agent model offers unique insights into team interactions and workflows, making it ideal for scenarios like software development and collaborative projects.

Limitations:

- **Learning Curve:** Due to its complexity, MetaGPT requires a deep understanding of both multi-agent systems and role-based collaboration to use effectively.
- **Documentation:** Similar to other emerging frameworks, the documentation for MetaGPT may not yet be comprehensive enough for beginners.

MetaGPT vs. AutoGen: While AutoGen mimics team dynamics, it can be less efficient than MetaGPT, where agents autonomously make communication decisions based on their roles. AutoGen’s hardcoded conditions for agent communication limit adaptability, while MetaGPT’s agents dynamically coordinate without predefined paths, allowing for greater flexibility in team-based workflows.

MetaGPT vs. CrewAI: While CrewAI focuses on simple agent processes, MetaGPT models entire team dynamics, allowing for more complex, adaptive interactions between agents. Additionally, CrewAI’s reliance on LangChain’s toolkits, which were not originally designed for function-calling models, leaves it vulnerable to errors, whereas MetaGPT’s decentralized agent approach ensures greater robustness and adaptability in dynamic, large-scale environments.

2.2.6. AutoGPT

AutoGPT is an experimental open-source application that showcases the capabilities of advanced Large Language Models (LLMs), such as GPT-4, in autonomously performing complex tasks without human intervention. Developed to push the boundaries of what is possible with AI, **AutoGPT can self-prompt and recursively break down tasks into subtasks**, tackling each component to achieve the user’s overarching goal.

AutoGPT operates by chaining together ”thoughts” generated by the LLM, allowing it to plan, execute, and learn from actions in a loop until the desired outcome is achieved. It integrates internet access for data gathering, long-term and short-term memory management, and can even write and execute code to perform specific tasks.

AutoGPT can be considered a type of sequential agent system, but it is not traditionally classified as a multi-agent system. Here is why:

- No multiple agents: MAS frameworks typically involve multiple agents that work in parallel or interact with each other to solve problems, with each agent playing a specialized role or having distinct behaviors. AutoGPT, in contrast, functions as a single-agent system, even though it may perform a variety of tasks that a MAS could handle.

- **No parallelism:** In MAS systems, agents often work in parallel, communicating and collaborating to solve complex tasks. AutoGPT lacks this parallelism and instead relies on a recursive, sequential problem-solving approach.

Some of the standout features of AutoGPT include:

- **Autonomous Task Management:** Capable of self-prompting and decomposing complex goals into manageable subtasks without human guidance.
- **Memory and Context Handling:** Utilizes both short-term and long-term memory to maintain context over extended interactions, enhancing its ability to perform multi-step tasks.
- **Internet and File System Access:** Can search the web for information and interact with the file system to read or write files as needed.
- **Code Generation and Execution:** Able to generate, read, and execute code, allowing it to perform a wide range of computational tasks.

Strengths:

- **Autonomy:** Excels in performing tasks end-to-end without human intervention, making it a powerful tool for automating complex workflows.
- **Memory and Context Management:** Efficiently manages context over long sessions, which is advantageous for applications requiring persistent state and long-term planning.
- **Flexibility:** Does not rely on predefined prompts or visual builders, allowing developers to define tasks and workflows directly in code.

Limitations:

- **Resource Intensive:** Running AutoGPT with powerful LLMs like GPT-4 can be computationally expensive and slow, which may limit its practicality for some applications.

- **Reliability Concerns:** As an experimental project, it may sometimes produce unexpected results or encounter issues like getting stuck in loops without achieving the desired outcome.
- **Lack of Fine-Grained Control:** While autonomous, it may lack the ability to be finely tuned for specific tasks compared to more specialized frameworks.

Best For: Developers and researchers interested in exploring the capabilities of autonomous agents powered by LLMs, particularly in automating complex, multi-step tasks without human intervention.

AutoGPT vs. AutoGen: In contrast, AutoGPT operates as a **single autonomous agent** that self-prompts and recursively tackles tasks. While both frameworks aim to leverage LLMs for task automation, AutoGPT emphasizes full autonomy without the need for inter-agent communication. AutoGen might offer more structure in multi-agent collaboration, whereas AutoGPT provides a more open-ended exploration of autonomous task completion.

2.3. Microchain

This framework is extremely underrated and not very actively maintained. However, its beauty lies in its extreme **simplicity**. Working with chain of thought is also extremely simple. This is definitely not production ready and seems to be more of a side project. I would love to see it getting some more love in the open source community.

2.4. Task-Specific MAS

2.4.1. Mindsearch:

MindSearch leverages a multi-agent system to mimic human cognitive processes, combining the strengths of search engines and Large Language Models (LLMs) to deliver accurate and comprehensive responses to user queries.

MindSearch addresses these challenges through a multi-agent framework composed of two main components:

- **WebPlanner:** The WebPlanner acts as the high-level orchestrator of the information retrieval process. It decomposes the user’s query into smaller, manageable sub-queries, modeling this decomposition as a dynamic graph. This graph represents the query and its sub-components as nodes, enabling a structured approach to information seeking. The WebPlanner uses a directed acyclic graph (DAG) to organize and process sub-queries, ensuring a logical flow of information retrieval steps.
- **WebSearcher:** WebSearcher executes the sub-queries generated by the WebPlanner. It performs hierarchical information retrieval using a coarse-to-fine strategy to manage large volumes of web content. The coarse level broadly generates multiple related queries to ensure wide coverage of the topic and executes these queries via search engines to gather a comprehensive set of initial results. The intermediate level aggregates and filters these results, removing duplicates and irrelevant entries, and selects the most valuable pages based on relevance. The fine level conducts detailed reading and summarization of selected pages, extracting key insights and integrating them into a coherent summary.

This strategy helps the model to gather information from numerous sources, offering more contextual search results than many private AI search engines, such as Perplexity.ai Pro and ChatGPT-Web. MindSearch also offers a variety of user interfaces such as React, Gradio, Streamlit, and Terminal to suit different needs.

This framework’s open-source nature invites further development and optimization by the research community. MindSearch sets a new standard for AI-driven search solutions, paving the way for more advanced and efficient information retrieval systems.

2.4.2. Haystack:

Haystack is an end-to-end LLM framework that enables the development of applications powered by Large Language Models (LLMs), Transformer models, vector search, and more. Whether the use case involves retrieval-augmented generation

(RAG), document search, question answering, or answer generation, Haystack can orchestrate state-of-the-art embedding models and LLMs into pipelines to build complete NLP applications. Developers can utilize **deepset Studio** to visually design and export Haystack pipeline architectures as YAML files or Python code.

While Haystack itself is not traditionally categorized as a Multi-Agent System (MAS) framework, it was originally developed for semantic search and document retrieval. However, with the recent introduction of "agents," Haystack has extended its capabilities in a manner that makes it more comparable to MAS frameworks, particularly in terms of task delegation and dynamic problem-solving. Inspired by the MRKL and ReAct frameworks, these Haystack agents enable LLMs to autonomously make decisions and perform actions to resolve complex queries. For instance, an agent may decide to search the web or query internal company files depending on the user's input. This dynamic interaction with tools and data sources mirrors the functionalities seen in some MAS frameworks, where multiple agents work together to achieve a common goal.

Key Features:

- **Agent-Led Information Retrieval:** Haystack agents autonomously determine which resources to query and how to present the retrieved information based on user input.
- **MRKL and ReAct Model Influences:** These frameworks inspire the agents' decision-making processes, allowing them to dynamically select tools and actions.
- **Extensibility:** Although not a fully-fledged MAS framework, Haystack is extensible enough to simulate multi-agent dynamics in information-driven tasks.

Strengths:

- **Advanced NLP Capabilities:** Haystack's integration with LLMs provides agents with sophisticated natural language understanding, making it highly effective for tasks such as information retrieval, question answering, and knowledge extraction.

- **Dynamic Problem Solving:** The agents can flexibly decide which tools to use, mimicking some aspects of agent coordination found in MAS frameworks.

Weaknesses:

- **Task-Specific:** While Haystack agents can manage search-related tasks, the framework does not fully support the development of general-purpose multi-agent systems like MADkit or AutoGen.
- **Not Truly Decentralized:** In contrast to MADkit, Haystack’s agents are not fully autonomous; they rely on centralized decision-making provided by LLMs.

Haystack is particularly well-suited for applications that involve exploring and searching through your own data. It is widely regarded for its stability and comprehensive documentation, making it an excellent choice for projects requiring question answering or semantic search capabilities.

2.4.3. OpenDevin

OpenDevin is the community version of Devin, a cutting-edge autonomous AI software engineer. It is designed to be a self-sufficient software developer, capable of planning, coding, testing, debugging, and deploying software without human intervention. OpenDevin’s mission is to replicate and exceed the abilities of the original Devin system. It focuses on automating all aspects of the software development lifecycle through intelligent task delegation, enabling complete end-to-end software creation autonomously.

Strengths:

- **Autonomy:** Fully automates the software development lifecycle, providing minimal need for human intervention.
- **Robustness:** Designed to handle large-scale software projects with scalability and adaptability.

Limitations:

- **Emerging Community Ecosystem:** As a community-driven project, it may lack the resources of commercial alternatives.
- **Complexity:** High learning curve due to its full-stack automation capabilities.

There are other frameworks, such as Devika, that aim to emulate Devin, but they still have considerable progress to make.

2.4.4. ChatDEV

ChatDEV operates as a virtual software company. Like MetaGPT, it assigns agents specific roles such as Chief Executive Officer (CEO), Chief Product Officer (CPO), Chief Technology Officer (CTO), programmer, reviewer, tester, and art designer. These agents collaborate to simulate a multi-agent organizational structure, allowing the team to coordinate in developing, testing, and deploying software projects. It is a structured MAS focused on predefined roles for each agent, making it particularly suited for simulating complex corporate environments in software development.

Strengths:

- **Organizational Simulation:** Excellent for simulating real-world team dynamics in software development.
- **Coordination:** Built-in support for coordination between multiple agents with specific roles.

Limitations:

- **Task-Specific:** Primarily focused on software development, limiting its use in other domains.
- **Rigid Role Assignment:** Agents are restricted to their defined roles, limiting flexibility in task execution.

2.4.5. Google AlphaCode

Google AlphaCode is a machine learning model designed to autonomously write code. It focuses on solving competitive programming problems, using LLMs to generate code that meets specific problem requirements. While AlphaCode doesn't operate like a traditional multi-agent system, it acts in a task-specific capacity for software code generation, focusing on efficient problem-solving.

Strengths:

- **High Accuracy:** Able to solve a wide range of programming challenges with minimal errors.
- **Rapid Code Generation:** Can produce code faster than traditional human coders in competitive settings.

Limitations:

- **Limited Scope:** Task-specific to competitive programming, not a general-purpose software development solution.
- **Non-Collaborative:** Doesn't use a team of agents working in unison but operates more like a solo agent.

Comparison (AlphaCode vs OpenDevin vs ChatDEV)

- **Task-Specific Nature:** While all three systems are task-specific, OpenDevin focuses on full-scale software development, ChatDEV simulates a corporate development environment, and AlphaCode specializes in competitive programming.
- **Autonomy:** OpenDevin and AlphaCode provide a high level of autonomy, with OpenDevin handling the entire software lifecycle, whereas ChatDEV delegates tasks to various agents with role-specific responsibilities.

- **Flexibility:** OpenDevin is more flexible as it handles all development tasks, while ChatDEV operates within a fixed corporate structure. AlphaCode, on the other hand, is restricted to coding problems.
- **Agent Collaboration:** ChatDEV excels in simulating multi-agent collaboration, while OpenDevin and AlphaCode are more autonomous, with less emphasis on multiple agents interacting.

3. Conclusion

For highly customized search experiences requiring precise control over agent interactions, LangGraph is an excellent choice. It excels in building sophisticated, stateful multi-agent systems that manage complex, cyclical workflows. In contrast, AutoGen is ideal for developers looking to create conversational AI applications with ease of use in mind. Its user-friendly interface and strong community support make it perfect for rapidly building, testing, and deploying multi-agent solutions. If you need similar ease of use as AutoGen but with more granular control over the system's inner workings, TaskWeaver offers the customization and control to fine-tune each aspect of the workflow.